

PyPy 1.0 and beyond



PyCon 07, Addison TX

Michael Hudson mwh@python.net
Heinrich-Heine-Universität Düsseldorf

What you're in for in the next 45 mins



- Quick intro and motivation
- Quick overview of architecture
- Some demos of features unique to PyPy, including the JIT
- A little talk about what the future holds

What is PyPy?



- PyPy is:
 - An implementation of Python, and a very flexible compiler framework (with some features that are especially useful for implementing interpreters)
 - An open source project (MIT license)
 - A STREP (“Specific Targeted REsearch Project”), partially funded by the EU
 - A lot of fun!

Status



- We can (still :-) produce a binary that looks very much like CPython to the user
- It's still a bit slower, depending on options and what you're doing
- Can also produce binaries that are more capable than CPython – with stackless-style coroutines, with logic variables, ...
- Can also produce binary for CLR (i.e. .NET)

Motivation



- PyPy grew out of a desire to modify/extend the *implementation* of Python, for example to:
 - increase performance (psyco-style JIT compilation, better garbage collectors)
 - add expressiveness (stackless-style coroutines, logic programming)
 - ease porting (to new platforms like the JVM or CLI or to low memory situations)

Lofty goals, but first...



- CPython is a fine implementation of Python but:
 - it's written in C, which makes porting to, for example, the CLI hard
 - while psyco and stackless exist, they are very hard to maintain as Python evolves
 - some implementation decisions are very hard to change (e.g. refcounting)

Enter the PyPy platform



Specification of the Python language

Compiler Framework

Python
running on JVM

Python
with JIT

Python for an
embedded device

Python with
transactional memory

Python just the way
you like it

How do you specify the Python language?



- The way we did it was to write an interpreter for Python in *RPython* – a subset of Python that is amenable to analysis
- This lets us write unit tests for our specification/implementation that run on top of CPython
- Can also test entire specification/implementation in same way

The “What is RPython?” question



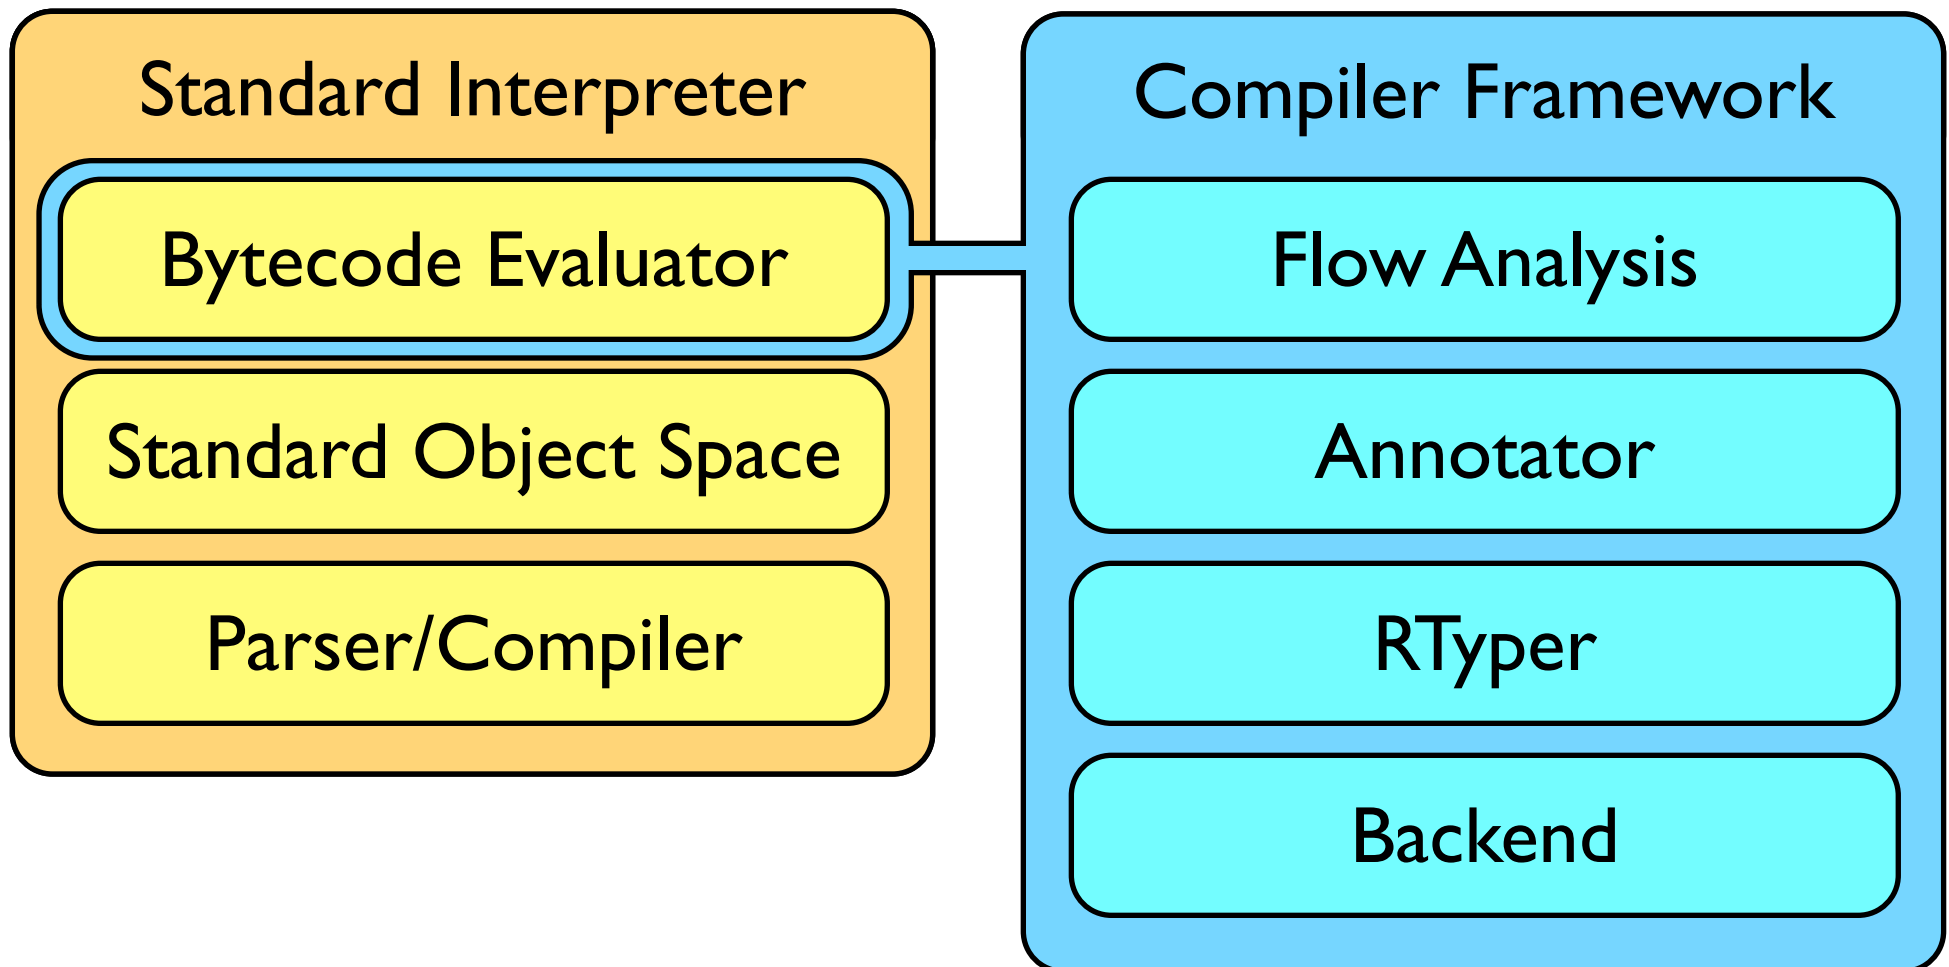
- Restricted Python, or RPython, first and foremost *is* Python
- It is a subset of Python that is static enough – *after initialization code has run* – for our analysis tools to cope with
- Somewhat Java-like – classes, methods, no pointers, no operator overloading – but at least has function pointers

Translation Aspects



- Our Python implementation is very high level
- One of our Big Goals is to produce our customized Python implementations without compromising on this point
- We do this by weaving in so-called ‘translation aspects’ during the compilation process

1,000 ft view



The Annotator



- Type annotation associates variables with information about which values they can take at run time
- An unusual feature of PyPy's approach is that the annotator works on live objects which means it never sees initialization code, so that can use `exec` and other dynamic tricks

The Annotator



- Annotation starts at a given entry point and discovers as it proceeds which functions may be called by the input program
- Does not modify the graphs; end result is essentially a big dictionary
- Read “Compiling dynamic language implementations” on the web site for more than is on these slides

The RTyper



- The RTyper takes as input an annotated RPython program (e.g. our Python implementation)
- It reduces the abstraction level of the graphs towards that of the target platform
- This is where the magic of PyPy really starts to get going :-)

The RTyper



- Can target a C-ish, pointer-using language or an OO language like Java or Smalltalk with classes and instances
- Resulting graphs are not completely low-level: still assume automatic memory management for example

Reducing Abstraction



- Many high level operations apply to different types – the most extreme example probably being calling an object
- For example, calling a function is RTyped to a `direct_call` but calling a class is RTyped to a sequence of operations including allocating the memory for the instance and calling any `__init__` function

Further Transforms



- RTyping is followed by a sequence of further transforms, depending on target platform and options supplied:
 - GC transformer – inserts explicit memory management operations
 - Stackless transform – inserts bookkeeping and extra operations to allow use of coroutines, tasklets etc
 - Exception transform – makes exception handling explicit

The Backend(s)



- Maintained backends: C, LLVM, CLI/.NET, and JavaScript
- All proceed in two phases:
 - Traverse the forest of rtyped graphs, computing names for everything
 - Spit out the code

Status – what works



- The Standard Interpreter very nearly complete
- The compiler framework:
 - Produces standalone binaries
 - C, LLVM and CLI backends well supported, JVM in progress
 - JavaScript backend works, but not for all of PyPy (not really intended to, either)

Status – what works



- The C backend support “stackless” features – coroutines, tasklets, recursion only limited by RAM
- Can use OS threads with a simple “GIL-thread” model
- Our Python specification/implementation has remained free of all these implementation decisions!

What we're working on now



- The Just-In-Time compiler – early stages, works for a very simple language
- More home-grown GCs (e.g. a semispace copying collector) and more GCs for LLVM
- Logic programming – some working code, interface and integration in progress

What we're working on now



- “rctypes”, a uniform way of calling external functions based on the now-standard “ctypes” module for CPython
- CLI (.NET) and Smalltalk backends
- supporting stackless features in other backends

About the project



- Open source, of course (MIT license)
- Distributed – the 12 paid developers live in 6 countries, contributors from more
- Sprint driven development – focussed week long coding sessions every ~6 weeks
- Extreme Programming practices: pair programming, test-driven development

“We’re Hiring!”



- In the open source sense:

- Read documentation:

<http://codespeak.net/pypy/>

- Come hang out in #pypy on freenode, post to pypy-dev
- Some opportunities for paid work too.